

IL Code Instrumentation with RAIL

Open source project offers easy, object-oriented modification of existing assemblies



Imagine that you have downloaded an application for the .NET platform, and that for some reason you would like to know what it is accessing on the hard drive. In particular, you would like to know which files it is reading from and writing to. This can be a problem since you may not have access to the source code of the application, and using Ildasm to iterate through the whole code is out of the question. In this article we will show you how to perform code instrumentation with the Runtime Assembly Instrumentation Library (RAIL), which allows assemblies to be easily modified according to high-level specifications.

► The Runtime Assembly Instrumentation Library

RAIL fills the gap that exists between .NET's reflection mechanisms and Reflection.Emit. While it is possible to use reflection to collect information about assemblies and types, and with Reflection.Emit it is possible to generate new assemblies at runtime, RAIL allows the modification of existing assemblies in an easy and object-oriented way. Some of the available high-level modifications include the ability to alter references in the code to types, fields, properties, and methods; insert prologue and epilogue instructions in method bodies; replace access to fields by methods; and selectively copy parts of assemblies. Many of these functionalities benefit from the use of design patterns to iterate through the object structure and apply the changes.

► A Simple Example

Let's consider the example mentioned earlier: how to monitor all disk accesses of an appli-

cation. With RAIL it is easy to replace all references to a certain class with new references to another class. In practice, to solve the problem it is necessary to replace all the references to classes that provide access to disk with references to an intermediate class. In this class, a log entry is created each time a method is called and then the original invocation can take place. The intermediate class is serving as a proxy between the original references and the classes that read and write from disk.

generated with a call to `this.FileWriteLog.WriteLine(fs.Name, v.ToString());`.

The FileWriteLog code in Listing 2 implements a very simple XML log maker. The constructor creates the log file, writes the start-of-document line and inserts the first element entry. The Close() method writes the end of the first entry, closes the document, and flushes the buffer. The WriteLine(string filePath, string val) method locks the access to the XmlTextWriter object to assure the atomicity of each log entry.

"Security is not something that can be pushed off onto the network administrator; it is an integral aspect of an application's overall architecture"

To completely solve the monitoring problem it would be necessary to proxy all the text reader and writer classes and also their binary counterparts. To keep the example short we will do this only for the System.IO.BinaryWriter class, creating a FileBinaryWriter class (see Figure 1).

To avoid concurrency problems when writing to the log, a new file is created every time the **FileBinaryWriter constructor** is called. This is achieved by dynamically generating the name of the log file using a static variable (LogVersion) in FileBinaryWriter. This variable is incremented every time it is used.

When overriding the Write() methods, a call is made to the base method and a log entry is

FileBinaryWriter and FileWriteLog are compiled under the MyFile namespace into the MyFile.dll assembly.

Foo.exe (see Listing 3) is the test assembly for this example. It is an adaptation of an MSDN example on how to use BinaryWriter objects. This short application writes the numbers from 0 to 10 into the Test.data file using BinaryWriter. It also checks if Test.data does not exist before the execution of Foo.exe.

► Putting RAIL to Work

In this example, RAIL will be used to copy the FileBinaryWriter and FileWriteLog types into the Foo.exe assembly, and to

BY BRUNO CABRAL &

LUIS SILVA



HOME



CLIENT



SERVER





replace the calls to `BinaryWriter` with calls to the newly inserted `FileBinaryWriter`. Listing 4 describes how this is done. Let's examine the implementation.

When using RAIL, the first thing to do is create a representation of the assembly to instrument using the `RAIL.Reflect.RAssemblyDef` class. The `RAssemblyDef.LoadAssembly()` creates an `RAssemblyDef` object, which represents the `Foo.exe` assembly. The same kind of object must also be created for the `MyFile.dll` assembly (see lines 15 and 16 of Listing 4).

The next step is to obtain `RType` objects representing the `FileBinaryWriter` and `FileWriteLog` classes of `MyFile.Assembly`. These instances are `FileBinaryWriterType` and `FileWriteLogType`, respectively.

Using the `CopyType()` method of `RAIL.Reflect.RModuleDef`, two copies of the `RType` objects are made into the `FooAssembly` instance. From these copy operations, two `RTypeDef` objects are returned.

The `ReferenceReplacer` class allows you to change references from one type to another. It is based on the Visitor design pattern. Line 28 of Listing 4 creates a `ReferenceReplacer` object for changing all the references in `NewFileBinaryWriterType` from `FileWriteLogType` to `NewFileWriteLogType`. The changes are applied using the `Accept()` method. This concludes the copying of the new type into the `FooAssembly`. Note that this type of copying step is not really necessary for the example; an external reference to the `MyFile.dll` assembly would be enough. Nevertheless, it is a nicer way to demonstrate RAIL's capabilities.

The next step is to create an `RType` object to represent the `BinaryWriter` class reference in the `FooAssembly`. Then, the `ReferenceReplacer` instance is instructed to redirect all the

references from `OriginalBinaryWriter` to `New FileBinaryWriterType`.

This `ReferenceReplacer` is applied to the `RTypeDef` instance of `Foo.MyStream` on line 38.

Finally, `FooAssembly.SaveAssembly("Foo.exe")` saves the new version of the assembly one directory up in the file system.

Figure 2 shows a screenshot of `Ildasm`. On the left the original assembly is shown; on the right the modified version is presented. It is easy to see that all references to `[mscorlib] System.IO.BinaryWriter` have been replaced with references to `MyFile.FileBinaryWriter`.

When executing the original `Foo.exe` a file named "test.data" is produced. When executing the modified version, a file called "log0.xml" will also be created. This file contains a log of the accesses that `Foo.exe` has made to the hard drive.

Conclusion

Although RAIL is still very new, it already provides powerful features in terms of code instrumentation, including code walkthrough, reference replacement, type coping, and method prologues and epilogues. The RAIL team is currently working in v1.0 of the software, which should be available by the end of the first trimester of 2004. Among other things, it will support assemblies with multiple modules.

The objective of RAIL is to simplify the task of code instrumentation, making it possible to use it in areas such as aspect-oriented programming, runtime analysis tools, security verification, and software fault injection.

Resource

- RAIL is an open source project of the Dependable Systems Group of the University of Coimbra, is distributed under the Mozilla Public License and is available for download from: <http://rail.dei.uc.pt> ●

Bruno Cabral, Paulo Marques, and Luis Silva are researchers in the Dependable Systems Group of the University of Coimbra, Portugal.

► bcabral@dei.uc.pt
► pmarques@dei.uc.pt
► luis@dei.uc.pt

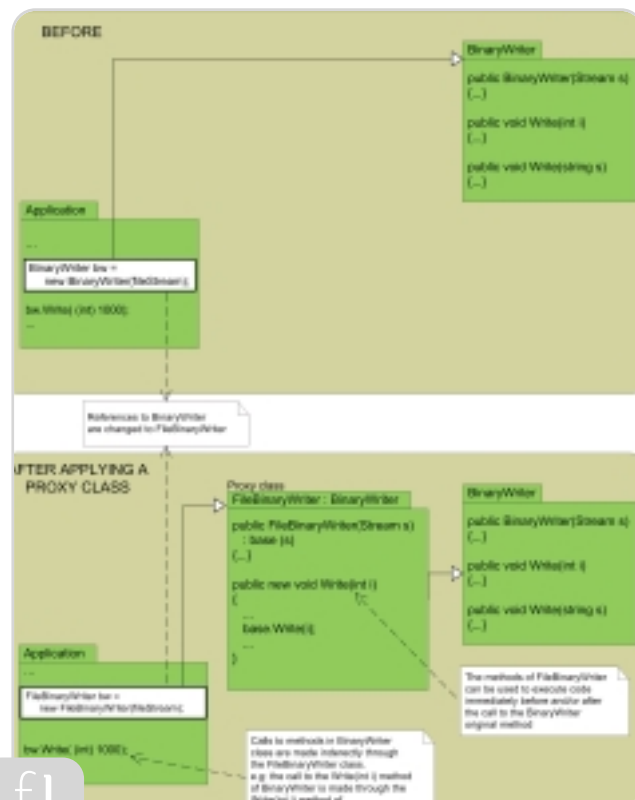


Figure 1: Creating a `FileBinaryWriter` class

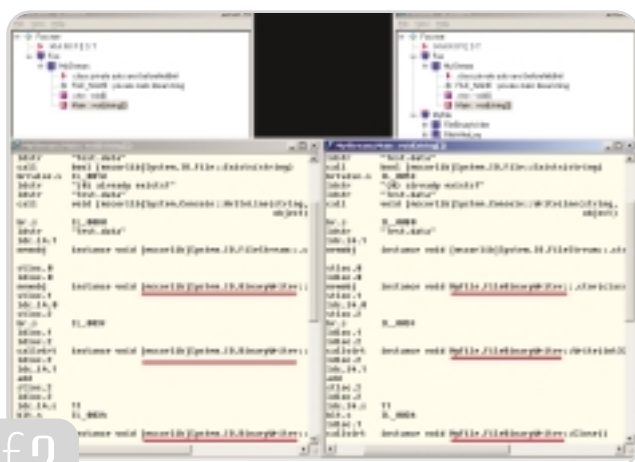


Figure 2: The Corporate wiki